

# Gender&Age\_Prediction

April 5, 2024

**Age and Gender Prediction.** In this project, we will be performing both classification and regression to predict both gender and age respectively.

```
[3]: # Ignore warnings for cleaner output
import warnings
warnings.filterwarnings('ignore')

# Data manipulation libraries
import pandas as pd
import numpy as np

# Operating system module
import os

# Visualization libraries
import matplotlib.pyplot as plt
import seaborn as sns

# Progress bar module
from tqdm.notebook import tqdm

# Tensorflow and Keras modules for building and training the model
import tensorflow as tf
from keras.preprocessing.image import load_img
from keras.models import Sequential, Model
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D, Input
```

**Load the dataset**

```
[ ]: # Mount Google Drive to access files
from google.colab import drive
drive.mount('/content/drive')

# Unzip the dataset file located in Google Drive
!unzip '/content/drive/MyDrive/Projects/AGE & GENDER/UTKFace.zip'

[5]: # Directory path where the dataset is located
BASE_DIR = "/content/UTKFace"
```

```
[6]: # List to store paths of images
image_paths = []

# List to store age labels
age_labels = []

# List to store gender labels
gender_labels = []
```

```
[7]: # Iterate over each file in the dataset directory
for filename in tqdm(os.listdir(BASE_DIR)):
    # Create the full path of the image file
    image_path = os.path.join(BASE_DIR, filename)

    # Split the filename to extract age and gender information
    temp = filename.split('_')

    # Extract age from the filename and convert it to integer
    age = int(temp[0])

    # Extract gender from the filename and convert it to integer
    gender = int(temp[1])

    # Append the image path to the list of image paths
    image_paths.append(image_path)

    # Append the age label to the list of age labels
    age_labels.append(age)

    # Append the gender label to the list of gender labels
    gender_labels.append(gender)
```

```
0%|          | 0/23708 [00:00<?, ?it/s]
```

```
[8]: print(f'Number of age_labels: {len(age_labels)}, Number of gender_labels: {len(gender_labels)}, Number of image_paths: {len(image_paths)}')
```

```
Number of age_labels: 23708, Number of gender_labels: 23708, Number of
image_paths: 23708
```

```
[9]: # Create a Pandas DataFrame to store image paths, age labels, and gender labels
df = pd.DataFrame()

# Assign image paths, age labels, and gender labels to DataFrame columns
df['image'] = image_paths
df['age'] = age_labels
df['gender'] = gender_labels
```

```
[10]: # Display the first few rows of the DataFrame to inspect its structure and
      ↪ contents
      df.head()
```

```
[10]:
```

	image	age	gender
0	/content/UTKFace/24_1_0_20170116222814643.jpg...	24	1
1	/content/UTKFace/25_0_1_20170116205335757.jpg...	25	0
2	/content/UTKFace/4_0_4_20161221195021183.jpg.c...	4	0
3	/content/UTKFace/28_0_0_20170117134849833.jpg...	28	0
4	/content/UTKFace/35_0_1_20170113152731945.jpg...	35	0

```
[11]: # Dictionary mapping gender labels to gender names
      gender_dict = {0: 'Male', 1: 'Female'}
```

### Exploratory Data Analysis (EDA)

```
[12]: # Import necessary module from PIL library
      from PIL import Image

      # Get the age and gender labels for the first image in the DataFrame
      age = df['age'][0]
      gender = df['gender'][0]

      # Open the first image using its path stored in the DataFrame
      img = Image.open(df['image'][0])

      # Set title with age and gender information
      plt.title(f'Age = {age} & Gender = {gender_dict[gender]}')

      # Turn off axis
      plt.axis('off')


      # Display the image
      plt.imshow(img);
```

Age = 24 & Gender = Female



```
[13]: # Plot a distribution plot for age labels using Seaborn  
sns.distplot(df['age'])
```

```
[13]: <Axes: xlabel='age', ylabel='Density'>
```



The distribution roughly follows a normal distribution that is slightly skewed to the right with a median of around 27 years. The majority are in between ages 25 to 30 years old. The range is from 0 to 120 years. There are some outliers at the higher end of the distribution.

```
[14]: # Set figure size
plt.figure(figsize=(20, 20))

# Select the first 50 rows from the DataFrame
files = df.iloc[0:50]

# Iterate over the selected rows
for index, file, age, gender in files.itertuples():
    # Create subplots in a grid of 10x5
    plt.subplot(10, 5, index+1)

    # Load image
    img = load_img(file)

    # Convert image to numpy array
    img = np.array(img)
```

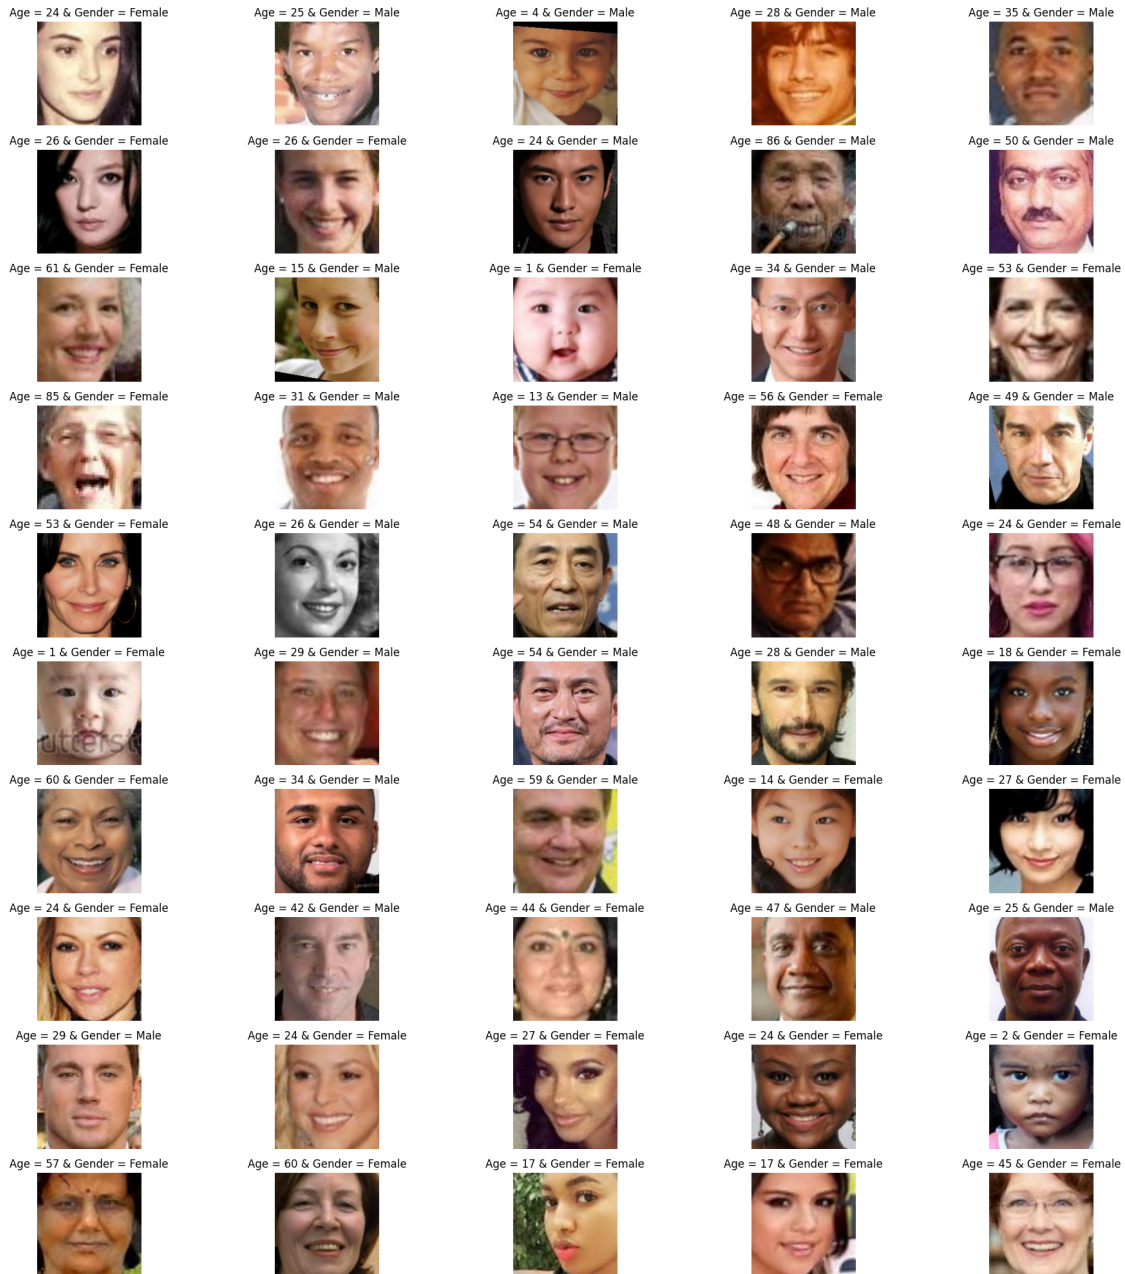
```
# Display the image
plt.imshow(img)

# Set title with age and gender information
plt.title(f'Age = {age} & Gender = {gender_dict[gender]}')

# Turn off axis
plt.axis('off')

# Adjust layout for better spacing
plt.tight_layout()

# Show the plot
plt.show()
```



## FEATURE EXTRACTION

```
[15]: # Define a function to extract features from images
def extract_features(images):
    # Initialize an empty list to store features
    features = []

    # Iterate over each image
    for image in tqdm(images):
```

```

    # Load image in grayscale mode and resize it to 128x128
    img = load_img(image, color_mode='grayscale')
    img = img.resize((128, 128), Image.LANCZOS)

    # Convert image to numpy array
    img = np.array(img)

    # Append the image to the features list
    features.append(img)

    # Convert the list of features to a numpy array
    features = np.array(features)

    # Reshape the features array to the required format for the model
    features = features.reshape(len(features), 128, 128, 1)

    return features

```

```

[16]: # Extract features from the images in the DataFrame
X = extract_features(df['image'])

```

```

0%|          | 0/23708 [00:00<?, ?it/s]

```

```

[17]: # Get the shape of the array representing extracted features
X.shape

```

```

[17]: (23708, 128, 128, 1)

```

```

[18]: # Normalize the pixel values of the images
X = X / 255.0

```

```

[19]: # Convert gender labels to numpy array
y_gender = np.array(df['gender'])

# Convert age labels to numpy array
y_age = np.array(df['age'])

```

```

[20]: # Define the input shape for the model
input_shape = (128, 128, 1)

```

```

[21]: from sklearn.model_selection import train_test_split

# Split the dataset into training and testing sets
X_train, X_test, y_gender_train, y_gender_test, y_age_train, y_age_test = \
    train_test_split(X, y_gender, y_age, test_size=0.2, random_state=42)

```

```

[22]: # Define the input layer
inputs = Input((input_shape))

```



```

# Convolutional layers
conv_1 = Conv2D(32, kernel_size=(3, 3), activation='relu')(inputs)
maxp_1 = MaxPooling2D(pool_size=(2, 2))(conv_1)
conv_2 = Conv2D(64, kernel_size=(3, 3), activation='relu')(maxp_1)
maxp_2 = MaxPooling2D(pool_size=(2, 2))(conv_2)
conv_3 = Conv2D(128, kernel_size=(3, 3), activation='relu')(maxp_2)
maxp_3 = MaxPooling2D(pool_size=(2, 2))(conv_3)
conv_4 = Conv2D(256, kernel_size=(3, 3), activation='relu')(maxp_3)
maxp_4 = MaxPooling2D(pool_size=(2, 2))(conv_4)

# Flatten layer
flatten = Flatten()(maxp_4)

# Fully connected layers
dense_1 = Dense(256, activation='relu')(flatten)
dense_2 = Dense(256, activation='relu')(flatten)

# Dropout layers
dropout_1 = Dropout(0.3)(dense_1)
dropout_2 = Dropout(0.3)(dense_2)

# Output layers for gender and age predictions
output_1 = Dense(1, activation='sigmoid', name='gender_out')(dropout_1)
output_2 = Dense(1, activation='relu', name='age_out')(dropout_2)

# Define the model with input and output layers
model = Model(inputs=[inputs], outputs=[output_1, output_2])

```

```

[23]: # Compile the model with appropriate loss functions, optimizer, and metrics
model.compile(loss=['binary_crossentropy', 'mae'],
              optimizer='adam',
              metrics=['accuracy', 'mae'])

```

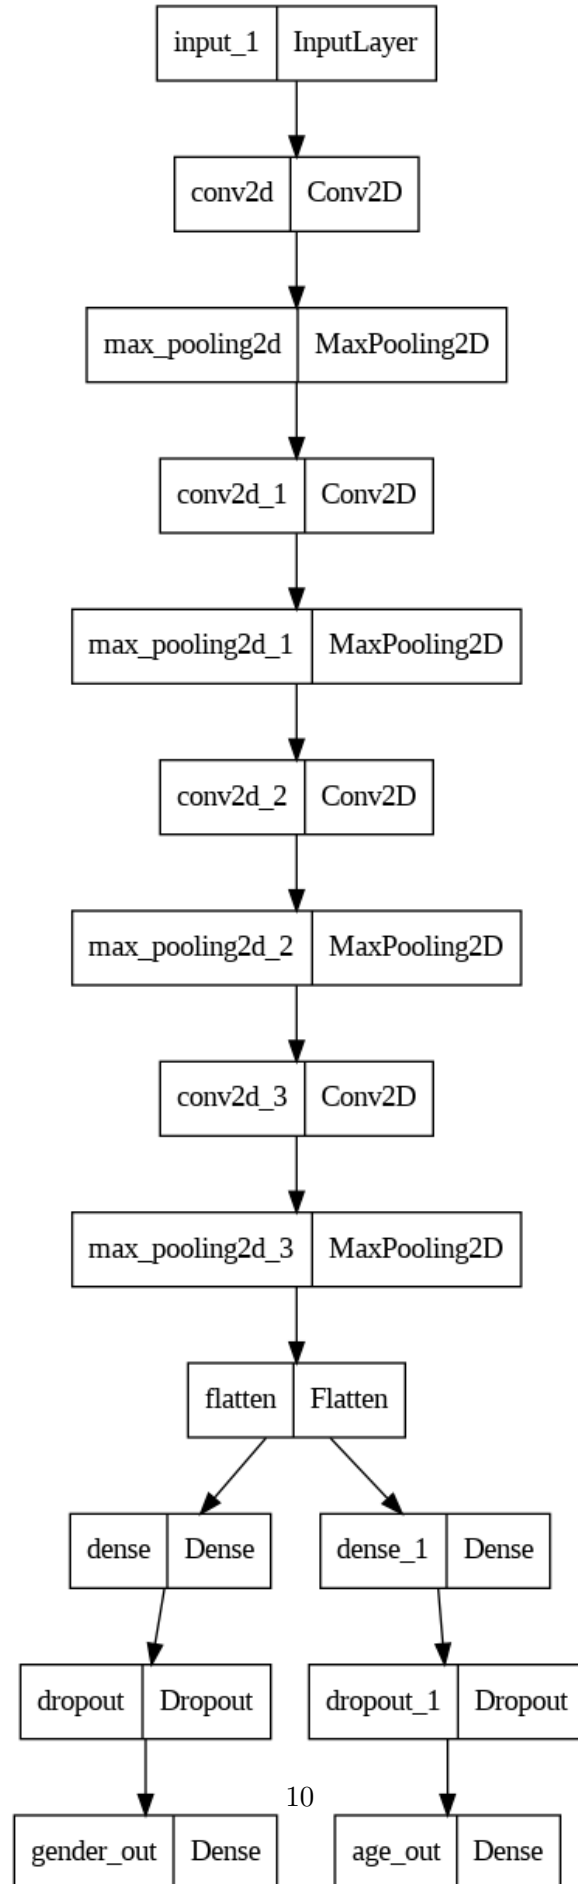
```

[24]: # Import the necessary module to plot the model
from tensorflow.keras.utils import plot_model

# Plot the model architecture
plot_model(model)

```

[24]:



```
[25]: # Train the model with the training data, using both gender and age labels as targets
      history = model.fit(x=X_train, y=[y_gender_train, y_age_train], batch_size=32, epochs=30, validation_split=0.2)
```

Epoch 1/30

```
475/475 [=====] - 22s 27ms/step - loss: 16.5882 -
gender_out_loss: 0.7081 - age_out_loss: 15.8801 - gender_out_accuracy: 0.5108 -
gender_out_mae: 0.4985 - age_out_accuracy: 0.0475 - age_out_mae: 15.8801 -
val_loss: 16.0832 - val_gender_out_loss: 0.6925 - val_age_out_loss: 15.3907 -
val_gender_out_accuracy: 0.5374 - val_gender_out_mae: 0.4996 -
val_age_out_accuracy: 0.0477 - val_age_out_mae: 15.3907
```

Epoch 2/30

```
475/475 [=====] - 10s 22ms/step - loss: 14.6125 -
gender_out_loss: 0.6225 - age_out_loss: 13.9900 - gender_out_accuracy: 0.6298 -
gender_out_mae: 0.4358 - age_out_accuracy: 0.0473 - age_out_mae: 13.9900 -
val_loss: 12.5151 - val_gender_out_loss: 0.4788 - val_age_out_loss: 12.0362 -
val_gender_out_accuracy: 0.7746 - val_gender_out_mae: 0.3353 -
val_age_out_accuracy: 0.0477 - val_age_out_mae: 12.0362
```

Epoch 3/30

```
475/475 [=====] - 11s 24ms/step - loss: 11.4935 -
gender_out_loss: 0.4650 - age_out_loss: 11.0284 - gender_out_accuracy: 0.7810 -
gender_out_mae: 0.3054 - age_out_accuracy: 0.0408 - age_out_mae: 11.0284 -
val_loss: 9.9371 - val_gender_out_loss: 0.3998 - val_age_out_loss: 9.5373 -
val_gender_out_accuracy: 0.8229 - val_gender_out_mae: 0.2766 -
val_age_out_accuracy: 0.0314 - val_age_out_mae: 9.5373
```

Epoch 4/30

```
475/475 [=====] - 11s 23ms/step - loss: 9.8988 -
gender_out_loss: 0.3972 - age_out_loss: 9.5016 - gender_out_accuracy: 0.8177 -
gender_out_mae: 0.2587 - age_out_accuracy: 0.0273 - age_out_mae: 9.5016 -
val_loss: 9.1260 - val_gender_out_loss: 0.3657 - val_age_out_loss: 8.7602 -
val_gender_out_accuracy: 0.8334 - val_gender_out_mae: 0.2336 -
val_age_out_accuracy: 0.0148 - val_age_out_mae: 8.7602
```

Epoch 5/30

```
475/475 [=====] - 11s 24ms/step - loss: 9.0532 -
gender_out_loss: 0.3633 - age_out_loss: 8.6899 - gender_out_accuracy: 0.8350 -
gender_out_mae: 0.2359 - age_out_accuracy: 0.0220 - age_out_mae: 8.6899 -
val_loss: 10.8331 - val_gender_out_loss: 0.3430 - val_age_out_loss: 10.4902 -
val_gender_out_accuracy: 0.8416 - val_gender_out_mae: 0.2240 -
val_age_out_accuracy: 0.0074 - val_age_out_mae: 10.4902
```

Epoch 6/30

```
475/475 [=====] - 11s 23ms/step - loss: 8.4760 -
gender_out_loss: 0.3348 - age_out_loss: 8.1412 - gender_out_accuracy: 0.8469 -
gender_out_mae: 0.2170 - age_out_accuracy: 0.0176 - age_out_mae: 8.1412 -
val_loss: 7.9622 - val_gender_out_loss: 0.3226 - val_age_out_loss: 7.6396 -
```

val\_gender\_out\_accuracy: 0.8527 - val\_gender\_out\_mae: 0.1999 -  
 val\_age\_out\_accuracy: 0.0098 - val\_age\_out\_mae: 7.6396  
 Epoch 7/30  
 475/475 [=====] - 11s 23ms/step - loss: 7.9371 -  
 gender\_out\_loss: 0.3101 - age\_out\_loss: 7.6270 - gender\_out\_accuracy: 0.8622 -  
 gender\_out\_mae: 0.1987 - age\_out\_accuracy: 0.0190 - age\_out\_mae: 7.6270 -  
 val\_loss: 7.6369 - val\_gender\_out\_loss: 0.3217 - val\_age\_out\_loss: 7.3152 -  
 val\_gender\_out\_accuracy: 0.8503 - val\_gender\_out\_mae: 0.1910 -  
 val\_age\_out\_accuracy: 0.0074 - val\_age\_out\_mae: 7.3152  
 Epoch 8/30  
 475/475 [=====] - 11s 23ms/step - loss: 7.5689 -  
 gender\_out\_loss: 0.2930 - age\_out\_loss: 7.2758 - gender\_out\_accuracy: 0.8704 -  
 gender\_out\_mae: 0.1867 - age\_out\_accuracy: 0.0173 - age\_out\_mae: 7.2758 -  
 val\_loss: 7.2670 - val\_gender\_out\_loss: 0.2904 - val\_age\_out\_loss: 6.9765 -  
 val\_gender\_out\_accuracy: 0.8730 - val\_gender\_out\_mae: 0.1776 -  
 val\_age\_out\_accuracy: 0.0074 - val\_age\_out\_mae: 6.9765  
 Epoch 9/30  
 475/475 [=====] - 10s 21ms/step - loss: 7.1821 -  
 gender\_out\_loss: 0.2732 - age\_out\_loss: 6.9089 - gender\_out\_accuracy: 0.8797 -  
 gender\_out\_mae: 0.1734 - age\_out\_accuracy: 0.0180 - age\_out\_mae: 6.9089 -  
 val\_loss: 7.4780 - val\_gender\_out\_loss: 0.2809 - val\_age\_out\_loss: 7.1970 -  
 val\_gender\_out\_accuracy: 0.8793 - val\_gender\_out\_mae: 0.1770 -  
 val\_age\_out\_accuracy: 0.0092 - val\_age\_out\_mae: 7.1970  
 Epoch 10/30  
 475/475 [=====] - 12s 24ms/step - loss: 6.9139 -  
 gender\_out\_loss: 0.2596 - age\_out\_loss: 6.6543 - gender\_out\_accuracy: 0.8851 -  
 gender\_out\_mae: 0.1629 - age\_out\_accuracy: 0.0167 - age\_out\_mae: 6.6543 -  
 val\_loss: 7.5915 - val\_gender\_out\_loss: 0.2817 - val\_age\_out\_loss: 7.3098 -  
 val\_gender\_out\_accuracy: 0.8819 - val\_gender\_out\_mae: 0.1785 -  
 val\_age\_out\_accuracy: 0.0132 - val\_age\_out\_mae: 7.3098  
 Epoch 11/30  
 475/475 [=====] - 11s 22ms/step - loss: 6.8010 -  
 gender\_out\_loss: 0.2516 - age\_out\_loss: 6.5494 - gender\_out\_accuracy: 0.8885 -  
 gender\_out\_mae: 0.1598 - age\_out\_accuracy: 0.0177 - age\_out\_mae: 6.5494 -  
 val\_loss: 7.2568 - val\_gender\_out\_loss: 0.2729 - val\_age\_out\_loss: 6.9839 -  
 val\_gender\_out\_accuracy: 0.8795 - val\_gender\_out\_mae: 0.1638 -  
 val\_age\_out\_accuracy: 0.0090 - val\_age\_out\_mae: 6.9839  
 Epoch 12/30  
 475/475 [=====] - 12s 25ms/step - loss: 6.4481 -  
 gender\_out\_loss: 0.2365 - age\_out\_loss: 6.2116 - gender\_out\_accuracy: 0.8961 -  
 gender\_out\_mae: 0.1499 - age\_out\_accuracy: 0.0182 - age\_out\_mae: 6.2116 -  
 val\_loss: 7.2439 - val\_gender\_out\_loss: 0.2732 - val\_age\_out\_loss: 6.9707 -  
 val\_gender\_out\_accuracy: 0.8793 - val\_gender\_out\_mae: 0.1722 -  
 val\_age\_out\_accuracy: 0.0055 - val\_age\_out\_mae: 6.9707  
 Epoch 13/30  
 475/475 [=====] - 11s 24ms/step - loss: 6.2267 -  
 gender\_out\_loss: 0.2243 - age\_out\_loss: 6.0024 - gender\_out\_accuracy: 0.9032 -  
 gender\_out\_mae: 0.1409 - age\_out\_accuracy: 0.0168 - age\_out\_mae: 6.0024 -

val\_loss: 7.9993 - val\_gender\_out\_loss: 0.2688 - val\_age\_out\_loss: 7.7305 -  
 val\_gender\_out\_accuracy: 0.8888 - val\_gender\_out\_mae: 0.1590 -  
 val\_age\_out\_accuracy: 0.0069 - val\_age\_out\_mae: 7.7305  
 Epoch 14/30  
 475/475 [=====] - 11s 23ms/step - loss: 5.9611 -  
 gender\_out\_loss: 0.2161 - age\_out\_loss: 5.7450 - gender\_out\_accuracy: 0.9053 -  
 gender\_out\_mae: 0.1377 - age\_out\_accuracy: 0.0173 - age\_out\_mae: 5.7450 -  
 val\_loss: 7.0924 - val\_gender\_out\_loss: 0.2684 - val\_age\_out\_loss: 6.8240 -  
 val\_gender\_out\_accuracy: 0.8898 - val\_gender\_out\_mae: 0.1437 -  
 val\_age\_out\_accuracy: 0.0090 - val\_age\_out\_mae: 6.8240  
 Epoch 15/30  
 475/475 [=====] - 11s 24ms/step - loss: 5.7146 -  
 gender\_out\_loss: 0.2016 - age\_out\_loss: 5.5130 - gender\_out\_accuracy: 0.9145 -  
 gender\_out\_mae: 0.1270 - age\_out\_accuracy: 0.0225 - age\_out\_mae: 5.5130 -  
 val\_loss: 7.1868 - val\_gender\_out\_loss: 0.2820 - val\_age\_out\_loss: 6.9048 -  
 val\_gender\_out\_accuracy: 0.8788 - val\_gender\_out\_mae: 0.1472 -  
 val\_age\_out\_accuracy: 0.0245 - val\_age\_out\_mae: 6.9048  
 Epoch 16/30  
 475/475 [=====] - 10s 22ms/step - loss: 5.4836 -  
 gender\_out\_loss: 0.1938 - age\_out\_loss: 5.2898 - gender\_out\_accuracy: 0.9170 -  
 gender\_out\_mae: 0.1229 - age\_out\_accuracy: 0.0432 - age\_out\_mae: 5.2898 -  
 val\_loss: 6.7811 - val\_gender\_out\_loss: 0.2661 - val\_age\_out\_loss: 6.5151 -  
 val\_gender\_out\_accuracy: 0.8962 - val\_gender\_out\_mae: 0.1505 -  
 val\_age\_out\_accuracy: 0.0414 - val\_age\_out\_mae: 6.5151  
 Epoch 17/30  
 475/475 [=====] - 12s 25ms/step - loss: 5.2717 -  
 gender\_out\_loss: 0.1844 - age\_out\_loss: 5.0873 - gender\_out\_accuracy: 0.9218 -  
 gender\_out\_mae: 0.1168 - age\_out\_accuracy: 0.0432 - age\_out\_mae: 5.0873 -  
 val\_loss: 7.1845 - val\_gender\_out\_loss: 0.2894 - val\_age\_out\_loss: 6.8951 -  
 val\_gender\_out\_accuracy: 0.8859 - val\_gender\_out\_mae: 0.1358 -  
 val\_age\_out\_accuracy: 0.0414 - val\_age\_out\_mae: 6.8951  
 Epoch 18/30  
 475/475 [=====] - 11s 23ms/step - loss: 5.1785 -  
 gender\_out\_loss: 0.1717 - age\_out\_loss: 5.0068 - gender\_out\_accuracy: 0.9268 -  
 gender\_out\_mae: 0.1103 - age\_out\_accuracy: 0.0447 - age\_out\_mae: 5.0068 -  
 val\_loss: 6.9792 - val\_gender\_out\_loss: 0.2791 - val\_age\_out\_loss: 6.7000 -  
 val\_gender\_out\_accuracy: 0.8890 - val\_gender\_out\_mae: 0.1419 -  
 val\_age\_out\_accuracy: 0.0456 - val\_age\_out\_mae: 6.7000  
 Epoch 19/30  
 475/475 [=====] - 11s 24ms/step - loss: 4.9589 -  
 gender\_out\_loss: 0.1633 - age\_out\_loss: 4.7956 - gender\_out\_accuracy: 0.9316 -  
 gender\_out\_mae: 0.1043 - age\_out\_accuracy: 0.0461 - age\_out\_mae: 4.7956 -  
 val\_loss: 7.1478 - val\_gender\_out\_loss: 0.3038 - val\_age\_out\_loss: 6.8439 -  
 val\_gender\_out\_accuracy: 0.8853 - val\_gender\_out\_mae: 0.1355 -  
 val\_age\_out\_accuracy: 0.0461 - val\_age\_out\_mae: 6.8439  
 Epoch 20/30  
 475/475 [=====] - 11s 23ms/step - loss: 4.8330 -  
 gender\_out\_loss: 0.1534 - age\_out\_loss: 4.6797 - gender\_out\_accuracy: 0.9340 -

gender\_out\_mae: 0.0986 - age\_out\_accuracy: 0.0461 - age\_out\_mae: 4.6797 -  
 val\_loss: 6.8887 - val\_gender\_out\_loss: 0.2969 - val\_age\_out\_loss: 6.5918 -  
 val\_gender\_out\_accuracy: 0.8822 - val\_gender\_out\_mae: 0.1361 -  
 val\_age\_out\_accuracy: 0.0469 - val\_age\_out\_mae: 6.5918  
 Epoch 21/30  
 475/475 [=====] - 11s 24ms/step - loss: 4.7012 -  
 gender\_out\_loss: 0.1450 - age\_out\_loss: 4.5562 - gender\_out\_accuracy: 0.9405 -  
 gender\_out\_mae: 0.0926 - age\_out\_accuracy: 0.0456 - age\_out\_mae: 4.5562 -  
 val\_loss: 6.9768 - val\_gender\_out\_loss: 0.3125 - val\_age\_out\_loss: 6.6643 -  
 val\_gender\_out\_accuracy: 0.8882 - val\_gender\_out\_mae: 0.1323 -  
 val\_age\_out\_accuracy: 0.0456 - val\_age\_out\_mae: 6.6643  
 Epoch 22/30  
 475/475 [=====] - 11s 24ms/step - loss: 4.6172 -  
 gender\_out\_loss: 0.1427 - age\_out\_loss: 4.4745 - gender\_out\_accuracy: 0.9392 -  
 gender\_out\_mae: 0.0933 - age\_out\_accuracy: 0.0453 - age\_out\_mae: 4.4745 -  
 val\_loss: 6.9775 - val\_gender\_out\_loss: 0.2953 - val\_age\_out\_loss: 6.6822 -  
 val\_gender\_out\_accuracy: 0.8911 - val\_gender\_out\_mae: 0.1301 -  
 val\_age\_out\_accuracy: 0.0472 - val\_age\_out\_mae: 6.6822  
 Epoch 23/30  
 475/475 [=====] - 11s 23ms/step - loss: 4.5018 -  
 gender\_out\_loss: 0.1282 - age\_out\_loss: 4.3736 - gender\_out\_accuracy: 0.9448 -  
 gender\_out\_mae: 0.0831 - age\_out\_accuracy: 0.0456 - age\_out\_mae: 4.3736 -  
 val\_loss: 7.2195 - val\_gender\_out\_loss: 0.3085 - val\_age\_out\_loss: 6.9110 -  
 val\_gender\_out\_accuracy: 0.8864 - val\_gender\_out\_mae: 0.1340 -  
 val\_age\_out\_accuracy: 0.0469 - val\_age\_out\_mae: 6.9110  
 Epoch 24/30  
 475/475 [=====] - 13s 28ms/step - loss: 4.3616 -  
 gender\_out\_loss: 0.1161 - age\_out\_loss: 4.2456 - gender\_out\_accuracy: 0.9515 -  
 gender\_out\_mae: 0.0756 - age\_out\_accuracy: 0.0459 - age\_out\_mae: 4.2456 -  
 val\_loss: 7.0612 - val\_gender\_out\_loss: 0.3299 - val\_age\_out\_loss: 6.7313 -  
 val\_gender\_out\_accuracy: 0.8898 - val\_gender\_out\_mae: 0.1339 -  
 val\_age\_out\_accuracy: 0.0453 - val\_age\_out\_mae: 6.7313  
 Epoch 25/30  
 475/475 [=====] - 11s 22ms/step - loss: 4.2698 -  
 gender\_out\_loss: 0.1152 - age\_out\_loss: 4.1546 - gender\_out\_accuracy: 0.9525 -  
 gender\_out\_mae: 0.0752 - age\_out\_accuracy: 0.0456 - age\_out\_mae: 4.1546 -  
 val\_loss: 6.9373 - val\_gender\_out\_loss: 0.3486 - val\_age\_out\_loss: 6.5887 -  
 val\_gender\_out\_accuracy: 0.8898 - val\_gender\_out\_mae: 0.1233 -  
 val\_age\_out\_accuracy: 0.0456 - val\_age\_out\_mae: 6.5887  
 Epoch 26/30  
 475/475 [=====] - 11s 24ms/step - loss: 4.2428 -  
 gender\_out\_loss: 0.1051 - age\_out\_loss: 4.1378 - gender\_out\_accuracy: 0.9568 -  
 gender\_out\_mae: 0.0683 - age\_out\_accuracy: 0.0455 - age\_out\_mae: 4.1378 -  
 val\_loss: 7.0968 - val\_gender\_out\_loss: 0.3610 - val\_age\_out\_loss: 6.7358 -  
 val\_gender\_out\_accuracy: 0.8801 - val\_gender\_out\_mae: 0.1343 -  
 val\_age\_out\_accuracy: 0.0472 - val\_age\_out\_mae: 6.7358  
 Epoch 27/30  
 475/475 [=====] - 11s 22ms/step - loss: 4.0552 -

gender\_out\_loss: 0.0985 - age\_out\_loss: 3.9567 - gender\_out\_accuracy: 0.9587 -  
gender\_out\_mae: 0.0652 - age\_out\_accuracy: 0.0453 - age\_out\_mae: 3.9567 -  
val\_loss: 7.1485 - val\_gender\_out\_loss: 0.3728 - val\_age\_out\_loss: 6.7757 -  
val\_gender\_out\_accuracy: 0.8911 - val\_gender\_out\_mae: 0.1248 -  
val\_age\_out\_accuracy: 0.0456 - val\_age\_out\_mae: 6.7757

Epoch 28/30

475/475 [=====] - 11s 23ms/step - loss: 4.0191 -  
gender\_out\_loss: 0.0928 - age\_out\_loss: 3.9263 - gender\_out\_accuracy: 0.9608 -  
gender\_out\_mae: 0.0608 - age\_out\_accuracy: 0.0463 - age\_out\_mae: 3.9263 -  
val\_loss: 7.1056 - val\_gender\_out\_loss: 0.4062 - val\_age\_out\_loss: 6.6993 -  
val\_gender\_out\_accuracy: 0.8846 - val\_gender\_out\_mae: 0.1272 -  
val\_age\_out\_accuracy: 0.0469 - val\_age\_out\_mae: 6.6993

Epoch 29/30

475/475 [=====] - 11s 23ms/step - loss: 3.9765 -  
gender\_out\_loss: 0.0875 - age\_out\_loss: 3.8890 - gender\_out\_accuracy: 0.9628 -  
gender\_out\_mae: 0.0574 - age\_out\_accuracy: 0.0453 - age\_out\_mae: 3.8890 -  
val\_loss: 7.0601 - val\_gender\_out\_loss: 0.3844 - val\_age\_out\_loss: 6.6756 -  
val\_gender\_out\_accuracy: 0.8896 - val\_gender\_out\_mae: 0.1258 -  
val\_age\_out\_accuracy: 0.0456 - val\_age\_out\_mae: 6.6756

Epoch 30/30

475/475 [=====] - 11s 23ms/step - loss: 3.9082 -  
gender\_out\_loss: 0.0863 - age\_out\_loss: 3.8219 - gender\_out\_accuracy: 0.9634 -  
gender\_out\_mae: 0.0554 - age\_out\_accuracy: 0.0462 - age\_out\_mae: 3.8219 -  
val\_loss: 7.0379 - val\_gender\_out\_loss: 0.3844 - val\_age\_out\_loss: 6.6535 -  
val\_gender\_out\_accuracy: 0.8940 - val\_gender\_out\_mae: 0.1241 -  
val\_age\_out\_accuracy: 0.0467 - val\_age\_out\_mae: 6.6535

## Plot Results


```
[26]: # Get training and validation accuracy for gender prediction
acc = history.history['gender_out_accuracy']
val_acc = history.history['val_gender_out_accuracy']
epochs = range(len(acc))

# Plot training and validation accuracy
plt.plot(epochs, acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'r', label='Validation Accuracy')
plt.title('Accuracy Graph For Gender Prediction')
plt.legend()
plt.show()


# Get training and validation loss for gender prediction
loss = history.history['loss']
val_loss = history.history['val_loss']

# Plot training and validation loss
plt.plot(epochs, loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Validation Loss')
```

```
plt.title('Loss Graph For Gender Prediction')  
plt.legend()  
plt.show()
```








```
[27]: # Get training and validation MAE for age prediction
loss = history.history['age_out_mae']
val_loss = history.history['val_age_out_mae']
epochs = range(len(loss))

# Plot training and validation MAE for age prediction
plt.plot(epochs, loss, 'b', label='Training MAE')
plt.plot(epochs, val_loss, 'r', label='Validation MAE')
plt.title('MAE Graph For Age Prediction')
plt.legend()
plt.show()
```



```
[28]: # Predictions on the test set
test_predictions = model.predict(X_test)

# Extracting gender predictions
gender_predictions = (test_predictions[0] > 0.5).astype(int).flatten()

# Calculating accuracy for gender prediction
from sklearn.metrics import accuracy_score
gender_accuracy = accuracy_score(y_gender_test, gender_predictions)

print("Gender Prediction Accuracy on Testing Set:", gender_accuracy*100,"%")

# Extracting age predictions
age_predictions = test_predictions[1].flatten()

# Calculating mean absolute error (MAE) for age prediction
from sklearn.metrics import mean_absolute_error
age_mae = mean_absolute_error(y_age_test, age_predictions)

print("Age Prediction MAE on Testing Set:", age_mae)
```

149/149 [=====] - 1s 7ms/step  
Gender Prediction Accuracy on Testing Set: 89.56136651202024 %  
Age Prediction MAE on Testing Set: 6.748970196750887

```
[29]: # Index of the image to be evaluated
image_index = 100

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 40  
1/1 [=====] - 0s 188ms/step  
Predicted Gender: Male Predicted Age: 37



```
[30]: # Index of the image to be evaluated
image_index = 3000

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

```
Original Gender: Female Original Age: 21
1/1 [=====] - 0s 21ms/step
Predicted Gender: Female Predicted Age: 28
```



```
[31]: # Index of the image to be evaluated
image_index = 10000

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:",
      y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Female Original Age: 12  
 1/1 [=====] - 0s 18ms/step  
 Predicted Gender: Female Predicted Age: 10



```
[32]: # Index of the image to be evaluated
image_index = 5000

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 32

1/1 [=====] - 0s 24ms/step

Predicted Gender: Male Predicted Age: 31



```
[33]: # Index of the image to be evaluated
image_index = 2000

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

```
Original Gender: Female Original Age: 36
1/1 [=====] - 0s 26ms/step
Predicted Gender: Female Predicted Age: 38
```



```
[37]: # Index of the image to be evaluated
image_index = 2300

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 2

1/1 [=====] - 0s 42ms/step

Predicted Gender: Male Predicted Age: 3





```
[35]: # Index of the image to be evaluated
image_index = 2900

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

```
Original Gender: Female Original Age: 8
1/1 [=====] - 0s 20ms/step
Predicted Gender: Female Predicted Age: 9
```



```
[36]: # Index of the image to be evaluated
image_index = 2709

# Print original gender and age labels
print("Original Gender:", gender_dict[y_gender[image_index]], "Original Age:", y_age[image_index])

# Predict gender and age labels using the model
pred = model.predict(X[image_index].reshape(1, 128, 128, 1))
pred_gender = gender_dict[round(pred[0][0][0])]
pred_age = round(pred[1][0][0])

# Print predicted gender and age labels
print("Predicted Gender:", pred_gender, "Predicted Age:", pred_age)

# Display the image
plt.axis('off')
plt.imshow(X[image_index].reshape(128, 128), cmap='gray');
```

Original Gender: Male Original Age: 21

1/1 [=====] - 0s 19ms/step

Predicted Gender: Male Predicted Age: 24

